



USI-CSCS Summer School 2013: Foundation of the Third Pillar of Science

William Sawyer, Themis Athanassiadou

CSCS-USI Summer School on Computer Simulations in Science and Engineering,
July 8-19 2013, Università Svizzera Italiana, Lugano, Switzerland


ETH
Eidgenössische Technische Hochschule Zürich
Swiss Federal Institute of Technology Zurich

CSCS
Swiss National Supercomputing Centre






Premise: 3 pillars of 21. century scientific method



Theory (since antiquity)



combined with experiment (since Galilei & Newton)



and simulation
(since Metropolis, Teller,
von Neumann, Fermi, ... 1940s)

**Excellence in Science requires excellence in all
three areas: theory, experiment, and simulations**

CSCS/USI Summer School

Goal: Introduce you to the *third pillar of science* through several key aspects of computer simulation:

- Basic concepts and tools
- “Best practices” for simulation software
- High performance computing (HPC)
 - * parallel processing (message passing, multithreading)
 - * analyzing performance
 - * libraries: utilizing third-party software
- A scientific use case: the Tsunami model

Organizers: CSCS, USI/ICS and the Foundations in Mathematics and Informatics for Computer Simulations in Science and Engineering (FoMICS) program



Summer School 2013 AGENDA

- July 8-9: HPC Software Engineering Best Practices
- July 10: Parallel programming introduction
- July 10 evening: summer school dinner
- July 11-12: Parallel programming with the Tsunami model
- July 13 or 14: hike (weather dependent)
- July 15-16: Parallel programming with Tsunami model continued
- July 16-18: Performance analysis
- July 17 evening: beach party (weather dependent)
- July 19: Numerical libraries

Beware: the summer school design is entirely new this year!



Some basic HPC concepts

- *Floating-point operation (Flop)*

✱ Flop per second (Flop/s)

- *Byte: 8 binary digits 0/1*


✱ Bytes per second (B/s, KB/s, MB/s, GB/s, ...)

- *Time to solution:* wall clock time (s.) to find answer
- *Energy to solution:* energy (Joules, KWh) to answer
- *Programming language:* language used describe calculation (generally at high level of abstraction)
- *Software engineering:* systematic approach for design/development/operation/maintenance of software

$$\begin{array}{r} 934,569.299814557 \\ \times \quad 52.827419489135904 \\ \hline = 49,370,884.442971624253823 \\ + \quad 4.20349729193958 \\ \hline = 49,370,888.64646892 \end{array}$$




Progress of computer power



ETH
Eidgenössische Technische Hochschule Zürich
Swiss Federal Institute of Technology Zurich

Koomey's Law (energy)




CSCS
Swiss National Supercomputing Centre

Computer performance and application performance increase $\sim 10^3$ every decade

~100 Kilowatts → ~5 Megawatts → 20-30 MW →
~1 Exaflop/s




1 GigaFlop/s
Cray YMP
8 processors



1.02 TeraFlop/s
Cray T3E
1'500 processors

1.35 PetaFlop/s
Cray XT5
150'000 processors



100 million or billion
processing cores (!)

1988	1998	2008	2018
First sustained GFlop/s Gordon Bell Prize 1988	First sustained TFlop/s Gordon Bell Prize 1998	First sustained PFlop/s Gordon Bell Prize 2008	Another 1,000x increase in sustained performance

Moore's gap: 10M processors is probably too optimistic, 100M processors with billions of threads more realistic
Opportunities for disruptive technologies, or maybe even a revolution?



Petaflop/s = 10^{15} 64-bit floating point operations / sec.

which takes more energy?

64-bit floating-point fused multiply add

or

moving three 64-bit operands 20 mm across the die

```
      934,569.299814557
x      52.827419489135904
-----
= 49,370,884.442971624253823
+      4.20349729193958
-----
= 49,370,888.64646892
```



20 mm

this takes over 3x the energy!

loading the data from off chip takes > 10x more yet


source: Steve Scott, Cray Inc.

moving data is expensive – exploiting data locality is critical to energy efficiency

If we care about energy consumption, we have to worry about these and other physical considerations of the computation – but where is the separation of concerns?

Current challenges

- CPU frequency not increasing
 - ➡ energy consumption
- Memory access times have lagged behind CPU performance

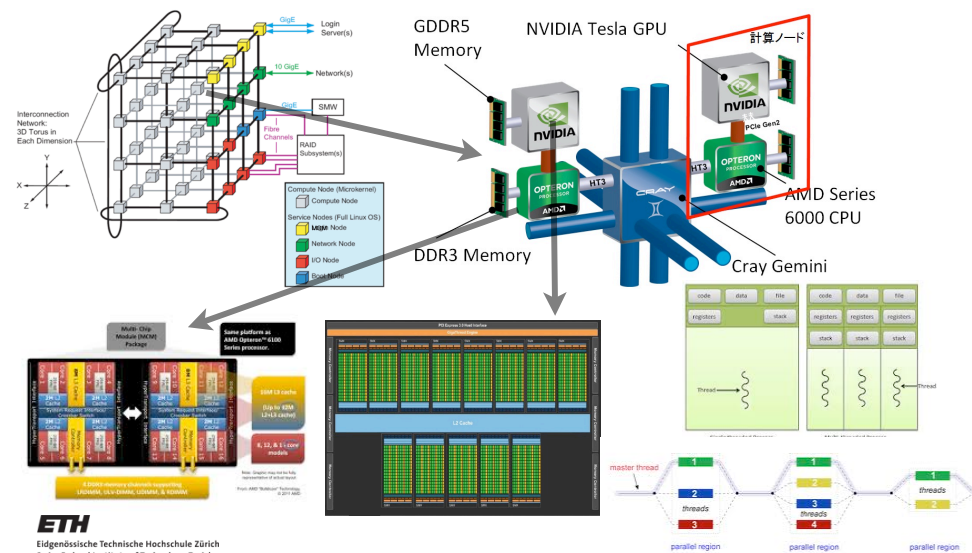


Solutions?

- Place multiple “cores” on same die within a “socket”
- Access data which is “nearby” (“caches”, local memory, ...)
- ➡ *requires the latest HPC programming techniques (and possibly some not invented yet...)*



State-of the-art HPC platform



Important CSCS access information

- See: *Programming environment for Practicals*
- Your account: `courseXX` (XX & pwd supplied separately)
- Access HPC platforms through CSCS front-end `ela.cscs.ch`
`ssh -Y First courseXX@ela.cscs.ch`
- Course platform: Cray XK7 (previous slide) `todi.cscs.ch`
`ssh -Y todi`
- For most of course: switch to the GNU programming env.
`module swap PrgEnv-cray PrgEnv-gnu`
- Compilers: C++: `cc` C: `cc` Fortran: `ftn`
- Allocating one node for interactive use:
`salloc --res=sschool -N 1 --time=01:00:00`
- Run a parallel job: `aprun -n 8 ./myprogram`

Course information on-line

- Software management site: <http://www.github.com>
 - * Repository: SummerSchool2013
 - * GitHub User: fomics
 - * Presentations: [SummerSchool2013/PRESENTATIONS](http://www.github.com/SummerSchool2013/PRESENTATIONS)
 - * Course information: [SummerSchool2013/wiki](http://www.github.com/SummerSchool2013/wiki)
- CSCS on-line (<http://www.cscs.ch>)
 - * User portal: <http://user.cscs.ch/>
- Course directory:
 - * `ela:/project/csstaff/courses/CSCS_USI_School`



Introduction to Software Engineering for Computer Simulations

Dr. William Sawyer

CSCS-USI Summer School on Computer Simulations in Science and Engineering,
July 8-19 2013, Università Svizzera Italiana, Lugano, Switzerland

Scientific simulation programming

In our experience:

- simulation software is written by scientists in the field
- these scientists are not trained in software engineering
- software often poorly designed/documented/implemented
- extensions often implemented by graduate students in the same field, again minimal training in software engineering. They adopt same style as the original code
- this development strategy is continued until (1) the software dies out, or (2) there is a major re-engineering effort

Poll: how many of you feel this applies to your project?



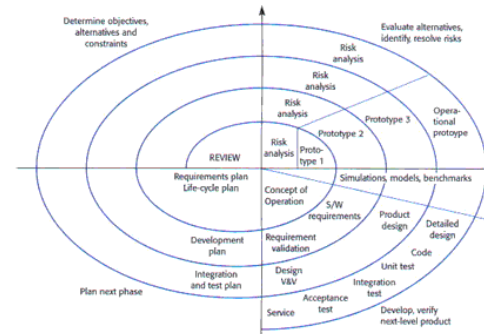
Software Engineering

[ACM Definition] *Software Engineering (SE) is concerned with developing and maintain software systems that:*

- *satisfy all the requirements that customers have defined*
- *behave reliably and efficiently*
- *are affordable*

SE introduces concepts like:

- *Requirements definition*
- *Development planning*
- *Prototyping*
- *Risk analysis*
- *Implementation / Validation*
- *Integration and acceptance*



ETH

Eidgenössische Technische Hochschule Zürich
Swiss Federal Institute of Technology Zurich

CSCS

Swiss National Supercomputing Centre



The case for Scientific Software Engineering

Traditionally, SE processes and techniques have not employed widely for scientific simulation software. *Why not?*

- Software is only a means to scientific results
- Perception: scientists and software engineers from 2 separate cultures
- Skilled scientists believe they can maintain overview of large software

Scientific Software Engineering can no longer be avoided:

- scientific models can be very large (> 1M lines of code)
- large numbers of concurrent developers
- programming difficulty of emerging hardware is increasing: *co-design* of scientists with hardware and software specialists is needed
- scientific computing can require huge investment (development and operating costs)



Two cultures?

Software engineer's
point of view:

Software engineer



I need your
requirements...

I need your
requirements
NOW ...

GIVE ME YOUR
REQUIREMENTS
...

Sigh ...

Sorry haven't quite
worked out what they
are...

Ooh, wouldn't it be
interesting if we
tried that? ...

Just have to work
out what's going on
here...


Sigh ...

Scientist



Software engineer's
point of view:

Scientist




Just write a
simple graph-
matching
program.

I need it by next
week.

EH?

Software engineer



What's a graph
matching program?

What sort of graphs?

How are they
matched?

How will it be used?

BY NEXT WEEK???

HELP!

ETH

Eidgenössische Technische Hochschule Zürich
Swiss Federal Institute of Technology Zurich

Source: Judith Segal

CSCS

Swiss National Supercomputing Centre



Useful links

Interestingly: little applicable web material available

- *Software Engineering for Scientific Computing* (Phil Colella, UC Berkeley)
The Seven Dwarfs
- *Scientists and software engineers: A Tale of Two Cultures* (Judith Segal, Open Univ.)
- *Computational Science and Engineering Department* (Chris Greenough, STFC)
- *Repeating Our Mistakes? Software Engineering for HPC* (David Bernholdt, ORNL):
“... long standing disconnect between computational science and software engineering”
- *Custom state-of-the-art scientific software* (Painter, Matthews, LANL):
“Writing code isn’t the problem. Understanding the problem is the problem.”



Conclusions of David Bernholdt

- There is a long-standing disconnect between computational science and software engineering
- More software engineering research is needed to support HPC and computational science
 - ➡ Currently spotty, many gaps, lacks critical mass
- Are we repeating our mistakes?
 - ➡ Yes
- Do we have to continue repeating them?
 - ➡ No



Goals for this component

The concept of this course stems from discussions with David Ham (Imperial College) at Infrastructure for European Network for Earth Science workshop
If software engineers cannot be employed, then at least SE-knowledgeable scientists!

- Get you acquainted with common SE tools used in scientific computing
- Promote “best practices” and recipes for good scientific software
 - *Clean programming style including documentation
 - *Software configuration management
 - *Validation and verification (automated)
- Explain legal issues relevant for scientific programmers (particularly grad students)
- Avoid boring business and project management aspects

Your input is needed to refine the course for future generations



Scenario I: simultaneous development

A large scientific code being developed by geographically distributed team

- individuals usually work on separate parts, but there is often overlap
- different components need to be “welded” together
- other developers not concerned about “details” of your implementation

Questions:

- How to make sure that changes to code base do not conflict?
- How does an individual develop/test without affecting others?
- How is the software merged into one?
- How are the individual components validated? the entire code?



Scenario 2: extensions to existing software

Your professor gives you the code from his dissertation and tells you to make extensions to it

- It's poorly written and documented
- It is monolithic, all low-level details visible
- It performs poorly, or and runs “sequentially” on one core

Questions:

- How do you re-engineer the code to “clean it up”?
 - ➡ Code needs: modularity, data encapsulation, proper documentation
 - ➡ Software configuration management
 - ➡ Validation
- How to track the record of all the improvements to the code?
- How to make it run more efficiently on the latest architectures?



Software Best Practices

The objectives of software development should be to make

- code simple to use by those who use it, but don't need to know the details
- your component easy to manage by those who come after you
- it as easy as possible to locate problems / bugs in the code
- it as easy as possible to integrate/merge development into a single system

Best practices aid toward these objectives

- Code broken up into modules with class hierarchies and clear interfaces; generally a single developer maintains one module
- Clear programming style: adherence to guidelines, expressive documentation, use of documentation system
- ...



Software Configuration Management

SCM is the *task of tracking and controlling changes in the software*

- Configuration Control: *implementing a controlled change process*
- Process Management: *adherence to organization's development process*
- Build Management: *managing the process and tools used to build software*
- Defect Tracking: *making sure every defect is traceable back to the source*
- **Data Provenance:** sometimes called “lineage” or “pedigree,” is the description of the origins of a piece of data and the process by which it arrived in a database
- **Revision Control**
 - ➡ *management of changes to documents, computer programs, and other collections of information*
 - ➡ *embedded into numerous document systems, e.g., Excel, Word, etc.*
 - ➡ *Version Control Systems (VCS) for large software projects: generally stand-alone applications*
 - ▶ *can be centralized: client-server, developers check into a central server*
 - ▶ *or distributed: each developer has a local repository, can share changes globally*

A few words about *data provenance*

Data provenance is really mentioned in SCM. Why not?

- In many ways it goes beyond SCM
- In most software fields, the software is the product, but
- In science, the resulting *data are the product*, i.e., the software is the means
- Scientific professionalism required data be *reproducible for years*, but
 - ▶ data depend on the computer architecture, program, algorithms, operating system, compiler brand and version
 - ▶ some fields (e.g. climatology) even require bit-for-bit reproducibility

Data provenance is a kind of metadata: it specifies the derivation history of a data product from its original sources. Its precise description depends on the field. In science, it ensures that data are reliable and reproducible, and therefore usable in other scientific work

Agenda Software Engineering


Today:

- Introduction to Software Configuration Management (David)
- C++ and its Best Practices in the Nutshell (Jean)
- Introduction to the Shallow Water Model “Tsunami” (Will)
- Using GIT for version control + LAB (John)

Tomorrow:

- Building code with Cmake + LAB (John)
- Verification techniques and manufactured solutions (David)
- Legal Issues with scientific code development (David)
- LAB: team development of Tsunami





Brief Introduction to the Doxygen Documentation System

Dr. William Sawyer, staff TU Munich

CSCS-USI Summer School on Computer Simulations in Science and Engineering,
July 8-19 2013, Università Svizzera Italiana, Lugano, Switzerland

What is doxygen?

Doxygen is a documentation system for C++, C, Java, Objective-C, Fortran and other languages. It helps in

- Generating on-line documentation or offline reference manual from documented source files
- Extracting the code structure and visualising the relations between various elements by means of include dependency graphs, inheritance diagrams, and collaboration diagrams, which are all generated automatically



Doxygen steps

- Create a configuration file
- Document the code
- Run doxygen




Creating a configuration file

- Doxygen determines settings from configuration file
- Doxywizard is a GUI front-end for configuring and running doxygen



Doxywizard



Ways to configure
and run doxygen

- Wizard
- Expert
- Load



Ways of documenting code

- Document blocks or lines
- Document members
- Structural commands
- Create lists



Documenting a block

Specify comment blocks:

- Special documentation blocks: comment blocks with additional markings
- Brief and detailed descriptions (see later)

C/C++ syntax

- Extra * or ! For Comment Blocks
e.g.

```
/**  
 * ... text ...  
 */
```

or

```
/*!  
 * ... text ...  
 */
```



Documenting lines

Specify comment lines:

- For comment lines: addition slash or an exclamation point

C/C++ syntax

```
///  
/// ... text ...  
///
```

or

```
//!  
//!... text ...  
//!
```

Brief description

Specify a brief description of **C/C++ syntax**
block

- `\brief` command

```
/*! \brief Brief description.  
 *      Brief description continued.  
 *  
 * Detailed description starts here.  
 */
```



Document members

- Documentation after member specification in source file
- Use ‘<’ to indicate that member is located in front of descriptor

C/C++ syntax

```
int var ; /*!< Description    */  
or Brief Description  
int var ; /*!< Brief Description
```



Creating lists

- Bulleted Lists: column-aligned minus sign ‘-’
- Numbered Lists: column-aligned minus sign followed by hash ‘-#’
- Nesting level is maintained according to column alignment
- HTML commands: can be used inside comment

C/C++ syntax

```
/*!  
 * List of events  
 * -Event1  
 * -#Subevent1  
 * -#Subevent2\n  
 * subevent2 cont.  
 * -#Subevent3  
 * -Event2  
 * -Event3  
 */
```



Structural commands

- Starts with ‘\’ or ‘@’
- Main commands
 - `\author` {list of authors}
 - `\brief` {brief description}
 - `\date` {date description}
 - `\file` [<name>]
 - `\fn` (function name) used if comment block is not placed
 - `\param` <parameter-name> {parameter description}
 - `\return` {return value description}



Visual enhancement commands

`\a` - Special font

`\b` - bold

`\c` - Typewriter Font

`\arg` - Simple bulleted list, not nested

`\e` - Italics

`\em` - Emphasize word and in Italics

