

# SWE – Anatomy of a Parallel Shallow Water Code

CSCS-FoMICS-USI Summer School on  
Computer Simulations in Science and Engineering

Michael Bader

July 8–19, 2013



# Teaching Parallel Programming Models ...

## Starting Point: Lecture on Parallel Programming

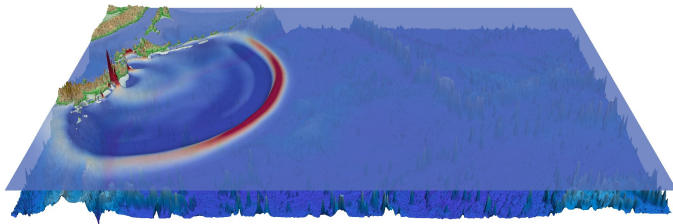
- classical approaches for shared & distributed memory: OpenMP and MPI
- “something more fancy” → GPU computing (CUDA, e.g.)
- motivating example to teach different models and compare their properties

## “Motivating Example”:

- not just Jacobi or Gauß-Seidel
- not the heat equation again ...
- inspired by a CFD code: “Nast” by Griebel et al.
- turned out to become shallow water equations
- **and is heavily used for summer schools ...**



# Towards Tsunami Simulation with SWE

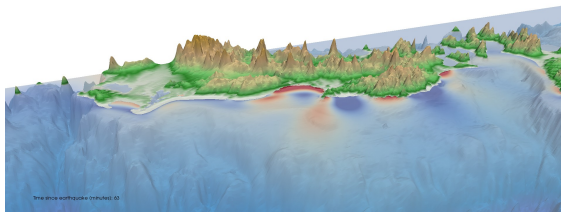


Time since earthquake (minutes): 78

## Shallow Water Code – Summary

- Finite Volume discretization on regular Cartesian grids  
→ simple numerics (but can be extended to state-of-the-art)
- patch-based approach with ghost cells for communication  
→ wide-spread design pattern for parallelization

# Towards Tsunami Simulation with SWE (2)

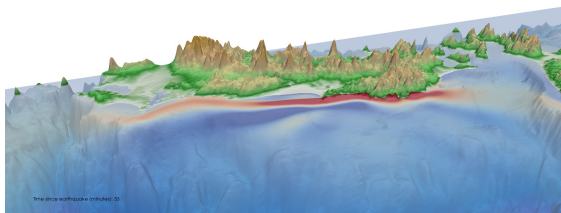


## Shallow Water Code – Bells & Whistles

- included augmented Riemann solvers (D. George, R. LeVeque)  
→ allows to simulate inundation
- developed towards hybrid parallel architectures  
→ now runs on GPU clusters

# Part I

## Model and Discretization

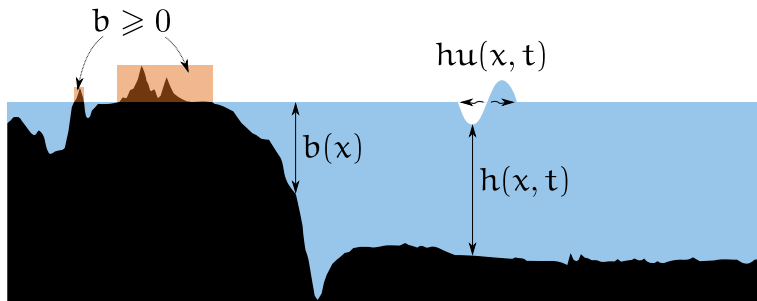


# Model: The Shallow Water Equations

Simplified setting (no friction, no viscosity, no coriolis forces, etc.):

$$\frac{\partial}{\partial t} \begin{bmatrix} h \\ hu \\ hv \end{bmatrix} + \frac{\partial}{\partial x} \begin{bmatrix} hu \\ hu^2 + \frac{1}{2}gh^2 \\ huv \end{bmatrix} + \frac{\partial}{\partial y} \begin{bmatrix} hv \\ huv \\ hv^2 + \frac{1}{2}gh^2 \end{bmatrix} = \begin{bmatrix} 0 \\ -\frac{\partial}{\partial x}(ghb) \\ -\frac{\partial}{\partial y}(ghb) \end{bmatrix}$$

**Quantities and unknowns:**



# Model: The Shallow Water Equations

Simplified setting (no friction, no viscosity, no coriolis forces, etc.):

$$\frac{\partial}{\partial t} \begin{bmatrix} h \\ hu \\ hv \end{bmatrix} + \frac{\partial}{\partial x} \begin{bmatrix} hu \\ hu^2 + \frac{1}{2}gh^2 \\ huv \end{bmatrix} + \frac{\partial}{\partial y} \begin{bmatrix} hv \\ huv \\ hv^2 + \frac{1}{2}gh^2 \end{bmatrix} = \begin{bmatrix} 0 \\ -\frac{\partial}{\partial x}(ghb) \\ -\frac{\partial}{\partial y}(ghb) \end{bmatrix}$$

**Write as generalized hyperbolic PDE:**

- 2D setting, three quantities:  $q = (q_1, q_2, q_3)^T = (h, hu, hv)^T$

$$\frac{\partial}{\partial t} q + \frac{\partial}{\partial x} F(q) + \frac{\partial}{\partial y} G(q) = S(q, x, y, t)$$

- with flux functions:

$$F(q) = \begin{bmatrix} q_2 \\ q_2^2/q_1 + \frac{1}{2}gq_1^2 \\ q_2q_3/q_1 \end{bmatrix} \quad G(q) = \begin{bmatrix} q_3 \\ q_2q_3/q_1 \\ q_3^2/q_1 + \frac{1}{2}gq_1^2 \end{bmatrix}$$

# Model: The Shallow Water Equations

Simplified setting (no friction, no viscosity, no coriolis forces, etc.):

$$\frac{\partial}{\partial t} \begin{bmatrix} h \\ hu \\ hv \end{bmatrix} + \frac{\partial}{\partial x} \begin{bmatrix} hu \\ hu^2 + \frac{1}{2}gh^2 \\ huv \end{bmatrix} + \frac{\partial}{\partial y} \begin{bmatrix} hv \\ huv \\ hv^2 + \frac{1}{2}gh^2 \end{bmatrix} = \begin{bmatrix} 0 \\ -\frac{\partial}{\partial x}(ghb) \\ -\frac{\partial}{\partial y}(ghb) \end{bmatrix}$$

## Derived from conservations laws:

- $h$  equation: conservation of mass
- equations for  $hu$  and  $hv$ : conservation of momentum
- $\frac{1}{2}gh^2$ : averaged hydrostatic pressure due to water column  $h$ , similar: bathymetry terms  $-\frac{\partial}{\partial x}(ghb)$  and  $-\frac{\partial}{\partial y}(ghb)$
- may also be derived by vertical averaging from the 3D incompressible Navier-Stokes equations



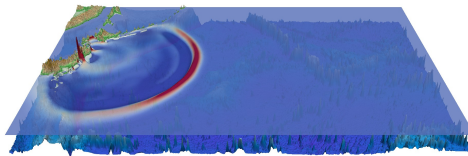
# Model: The Shallow Water Equations

Simplified setting (no friction, no viscosity, no coriolis forces, etc.):

$$\frac{\partial}{\partial t} \begin{bmatrix} h \\ hu \\ hv \end{bmatrix} + \frac{\partial}{\partial x} \begin{bmatrix} hu \\ hu^2 + \frac{1}{2}gh^2 \\ huv \end{bmatrix} + \frac{\partial}{\partial y} \begin{bmatrix} hv \\ huv \\ hv^2 + \frac{1}{2}gh^2 \end{bmatrix} = \begin{bmatrix} 0 \\ -\frac{\partial}{\partial x}(ghb) \\ -\frac{\partial}{\partial y}(ghb) \end{bmatrix}$$

## The ocean as “shallow water”??

- compare horizontal ( $\sim 1000$  km) to vertical ( $\sim 4$  km) length scale
- wave lengths large compared to water depth
- vertical flow may be neglected; movement of the “entire water column”



Time since earthquake (minutes): 75

# Finite Volume Discretisation

- discretise system of PDEs

$$\frac{\partial}{\partial t} q + \frac{\partial}{\partial x} F(q) + \frac{\partial}{\partial y} G(q) = S(t, x, y)$$

- results from integral equation:

$$\frac{\partial}{\partial t} \int_{t_n}^{t_{n+1}} \int_{\Omega} q \, d\omega \, dt + \int_{t_n}^{t_{n+1}} \int_{\partial\Omega} \vec{F}(q) \cdot \vec{n} \, ds \, dt = \dots$$

- use averaged quantities  $Q_{i,j}^{(n)}$  in finite volume elements  $\Omega_{ij}$ :

$$Q_{ij}(t) := \frac{1}{|\Omega_{ij}|} \int_{\Omega_{ij}} q \, d\omega \quad \rightsquigarrow \quad \frac{\partial}{\partial t} \int_{t_n}^{t_{n+1}} \int_{\Omega} q \, d\omega \, dt = |\Omega_{ij}| \left( Q_{i,j}^{(n+1)} - Q_{i,j}^{(n)} \right)$$

- What about the flux integral?

## Finite Volume Discretisation (2)

- flux integral on Cartesian grids:

$$\begin{aligned} \int_{t_n}^{t_{n+1}} \int_{\partial\Omega} \vec{F}(q) \cdot \vec{n} \, ds \, dt &= \int_{t_n}^{t_{n+1}} \int_{y_{j-\frac{1}{2}}}^{y_{j+\frac{1}{2}}} F(q(x_{i+\frac{1}{2}}, y, t)) - F(q(x_{i-\frac{1}{2}}, y, t)) \, dy \, dt \\ &\quad + \int_{t_n}^{t_{n+1}} \int_{x_{i-\frac{1}{2}}}^{x_{i+\frac{1}{2}}} G(q(x, y_{j+\frac{1}{2}}, t)) - G(q(x, y_{j-\frac{1}{2}}, t)) \, dx \, dt \end{aligned}$$

- leads to explicit time stepping scheme:

$$\begin{aligned} Q_{i,j}^{(n+1)} - Q_{i,j}^{(n)} &= \frac{\Delta t}{\Delta y} \left( F(q(x_{i+\frac{1}{2}}, y, t_n)) - F(q(x_{i-\frac{1}{2}}, y, t_n)) \right) \\ &\quad + \frac{\Delta t}{\Delta x} \left( G(q(x, y_{j+\frac{1}{2}}, t_n)) - G(q(x, y_{j-\frac{1}{2}}, t_n)) \right) \end{aligned}$$

- how to compute  $F_{i+\frac{1}{2},j}^{(n)} := F(q(x_{i+\frac{1}{2}}, y, t_n))$ ?

# Central and Upwind Fluxes

- define fluxes  $F_{i+\frac{1}{2},j}^{(n)}, G_{i,j+\frac{1}{2}}^{(n)}, \dots$  via 1D numerical flux function  $\mathcal{F}$ :

$$F_{i+\frac{1}{2}}^{(n)} = \mathcal{F}(Q_i^{(n)}, Q_{i+1}^{(n)}) \quad G_{j-\frac{1}{2}}^{(n)} = \mathcal{F}(Q_{j-1}^{(n)}, Q_j^{(n)})$$

- central flux:**

$$F_{i+\frac{1}{2}}^{(n)} = \mathcal{F}(Q_i^{(n)}, Q_{i+1}^{(n)}) := \frac{1}{2} \left( F(Q_i^{(n)}) + F(Q_{i+1}^{(n)}) \right)$$

leads to unstable methods for convective transport

- upwind flux** (here, for  $h$ -equation,  $F(h) = hu$ ):

$$F_{i+\frac{1}{2}}^{(n)} = \mathcal{F}(h_i^{(n)}, h_{i+1}^{(n)}) := \begin{cases} hu|_i & \text{if } u|_{i+\frac{1}{2}} > 0 \\ hu|_{i+1} & \text{if } u|_{i+\frac{1}{2}} < 0 \end{cases}$$

stable, but includes artificial diffusion

## (Local) Lax-Friedrichs Flux

- classical **Lax-Friedrichs method** uses as numerical flux:

$$F_{i+\frac{1}{2}}^{(n)} = \mathcal{F}(Q_i^{(n)}, Q_{i+1}^{(n)}) := \frac{1}{2} \left( F(Q_i^{(n)}) + F(Q_{i+1}^{(n)}) \right) - \frac{h}{2\tau} (Q_{i+1}^{(n)} - Q_i^{(n)})$$

- can be interpreted as central flux plus diffusion flux:

$$\frac{h}{2\tau} (Q_{i+1}^{(n)} - Q_i^{(n)}) = \frac{h^2}{2\tau} \cdot \frac{Q_{i+1}^{(n)} - Q_i^{(n)}}{h}$$

with diffusion coefficient  $\frac{h^2}{2\tau}$ , where  $c := \frac{h}{\tau}$  is a velocity  
("one grid cell per time step"  $\rightarrow$  cmp. CFL condition)

- idea of **local Lax-Friedrichs** method: use the actual wave speed

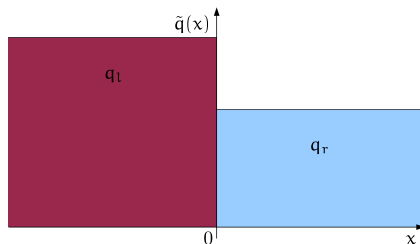
$$F_{i+\frac{1}{2}}^{(n)} := \frac{1}{2} \left( F(Q_i^{(n)}) + F(Q_{i+1}^{(n)}) \right) - \frac{a_{i+\frac{1}{2}}}{2} (Q_{i+1}^{(n)} - Q_i^{(n)})$$

# Riemann Problems

- solve **Riemann problem** to obtain solution  $q(x_{i+\frac{1}{2}}, y, t_n)$ , etc.:
- 1D treatment: solve shallow water equations with initial conditions

$$q(x_{i-\frac{1}{2}}, t_n) = \begin{cases} q_l = Q_{i-1}^{(n)} & \text{if } x < x_{i-\frac{1}{2}} \\ q_r = Q_i^{(n)} & \text{if } x > x_{i-\frac{1}{2}} \end{cases}$$

- solution: two (left or right) outgoing waves (shock or rarefaction)

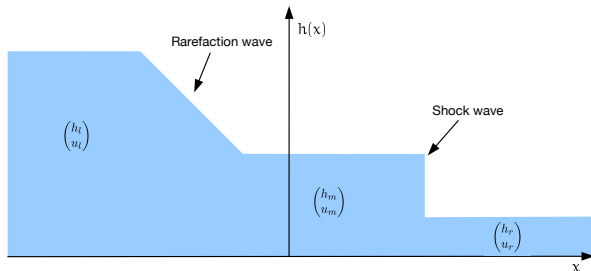


# Riemann Problems

- solve **Riemann problem** to obtain solution  $q(x_{i+\frac{1}{2}}, y, t_n)$ , etc.:
- 1D treatment: solve shallow water equations with initial conditions

$$q(x_{i-\frac{1}{2}}, t_n) = \begin{cases} q_l = Q_{i-1}^{(n)} & \text{if } x < x_{i-\frac{1}{2}} \\ q_r = Q_i^{(n)} & \text{if } x > x_{i-\frac{1}{2}} \end{cases}$$

- solution: two (left or right) outgoing waves (shock or rarefaction)

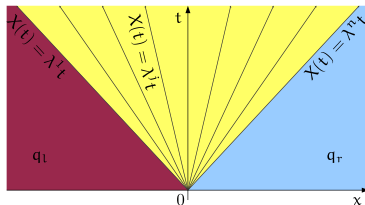


## Riemann Problems (2)

- wave propagation approach: split the jump into fluxes

$$F(Q_i) - F(Q_{i-1}) - \Delta x \psi_{i-\frac{1}{2}} = \sum_p \alpha_p r_p \equiv \sum_p Z_p \quad \alpha_p \in \mathbb{R}.$$

$r_p$  the eigenvector of the linearised problem,  
 $\psi_{i-\frac{1}{2}}$  a fix for the source term (bathymetry)



- implementation will compute **net updates**:

$$\mathcal{A}^+ \Delta Q_{i-1/2,j} = \sum_{p: \lambda_p > 0} Z_p \quad \mathcal{A}^- \Delta Q_{i-1/2,j} = \sum_{p: \lambda_p < 0} Z_p$$



# The F-Wave Solver

- use Roe eigenvalues  $\lambda_{1/2}^{\text{Roe}}$  to approximate the wave speeds:

$$\lambda_{1/2}^{\text{Roe}}(q_l, q_r) = u^{\text{Roe}}(q_l, q_r) \pm \sqrt{gh^{\text{Roe}}(q_l, q_r)}$$

- with  $h^{\text{Roe}}(q_l, q_r) = \frac{1}{2}(h_l + h_r)$  and  $u^{\text{Roe}}(q_l, q_r) = \frac{u_l\sqrt{h_l} + u_r\sqrt{h_r}}{\sqrt{h_l} + \sqrt{h_r}}$
- eigenvectors  $r_{1/2}^{\text{Roe}}$  for wave decomposition defined as

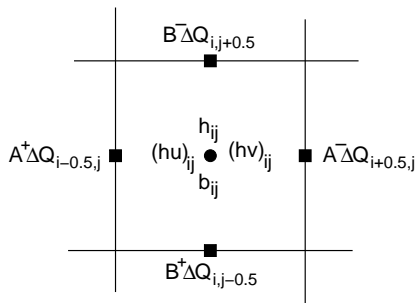
$$r_1^{\text{Roe}} = \begin{pmatrix} 1 \\ \lambda_1^{\text{Roe}} \end{pmatrix} \quad r_2^{\text{Roe}} = \begin{pmatrix} 1 \\ \lambda_2^{\text{Roe}} \end{pmatrix}$$

- leads to net updates (source terms still missing):

$$A^- \Delta Q := \sum_{p: \{\lambda_p^{\text{Roe}} < 0\}} \alpha_p r_p \quad A^+ \Delta Q := \sum_{p: \{\lambda_p^{\text{Roe}} > 0\}} \alpha_p r_p$$

- with  $\alpha_{1/2}$  computed from  $\begin{pmatrix} 1 \\ \lambda_1^{\text{Roe}} \end{pmatrix} \begin{pmatrix} 1 \\ \lambda_2^{\text{Roe}} \end{pmatrix} \begin{pmatrix} \alpha_1 \\ \alpha_2 \end{pmatrix} = F(Q_i) - F(Q_{i-1})$

# Finite Volume on Cartesian Grids



## Unknowns and Numerical Fluxes:

- (averaged) unknowns  $h$ ,  $hu$ ,  $hv$ , and  $b$  located in cell centers
- two sets of “net updates” or “numerical fluxes” per edge;  
here:  $\mathcal{A}^{+}\Delta Q_{i-1/2,j}$ ,  $\mathcal{B}^{-}\Delta Q_{i,j+1/2}$  (“wave propagation form”)

# Flux Form vs. Wave Propagation Form

- numerical scheme in flux form:

$$Q_{i,j}^{(n+1)} = Q_{i,j}^{(n)} - \frac{\Delta t}{\Delta x} \left( F_{i+\frac{1}{2},j}^{(n)} - F_{i-\frac{1}{2},j}^{(n)} \right) - \frac{\Delta t}{\Delta y} \left( G_{i,j+\frac{1}{2}}^{(n)} - G_{i,j-\frac{1}{2}}^{(n)} \right)$$

where  $F_{i+\frac{1}{2},j}^{(n)}$ ,  $G_{i,j+\frac{1}{2}}^{(n)}$ ,  $\dots$  approximate the flux functions  $F(q)$  and  $G(q)$  at the grid cell boundaries

- Wave propagation form:**

$$Q_{i,j}^{n+1} = Q_{i,j}^n - \frac{\Delta t}{\Delta x} \left( \mathcal{A}^+ \Delta Q_{i-1/2,j} + \mathcal{A}^- \Delta Q_{i+1/2,j}^n \right) - \frac{\Delta t}{\Delta y} \left( \mathcal{B}^+ \Delta Q_{i,j-1/2} + \mathcal{B}^- \Delta Q_{i,j+1/2}^n \right).$$

where  $\mathcal{A}^+ \Delta Q_{i-1/2,j}$ ,  $\mathcal{B}^- \Delta Q_{i,j+1/2}^n$ , etc. are **net updates**

- difference in implementation: compute one “flux term” or two “net updates” for each edge

# Time Stepping: Splitting or Not?

- With **Dimensional Splitting**:

$$Q_{i,j}^* = Q_{i,j}^n - \frac{\Delta t}{\Delta x} \left( \mathcal{A}^+ \Delta Q_{i-1/2,j}^n + \mathcal{A}^- \Delta Q_{i+1/2,j}^n \right)$$

$$Q_{i,j}^{n+1} = Q_{i,j}^* - \frac{\Delta t}{\Delta y} \left( \mathcal{B}^+ \Delta Q_{i,j-1/2}^* + \mathcal{B}^- \Delta Q_{i,j+1/2}^* \right).$$

two sequential “sweeps” of Riemann solves on horizontal vs. vertical edges

- vs. “un-split” method: (**currently used in SWE**)

$$Q_{i,j}^{n+1} = Q_{i,j}^n - \frac{\Delta t}{\Delta x} \left( \mathcal{A}^+ \Delta Q_{i-1/2,j}^n + \mathcal{A}^- \Delta Q_{i+1/2,j}^n \right) - \frac{\Delta t}{\Delta y} \left( \mathcal{B}^+ \Delta Q_{i,j-1/2}^n + \mathcal{B}^- \Delta Q_{i,j+1/2}^n \right).$$

allows to combine loops on horizontal and vertical edges

# Time Stepping

## CFL Condition:

- we only consider neighbour cells for a time step  
⇒ information must not travel faster than one cell per timestep!
- thus: timesteps need to consider characteristic wave speeds
- rule of thumb: wave speed depends on water depth,  $\lambda = \sqrt{gh}$
- in SWE: Riemann solvers will compute local wave speeds  
⇒ maximum-reduction necessary to find global time step

## Adaptive time step control forces sequential main loop:

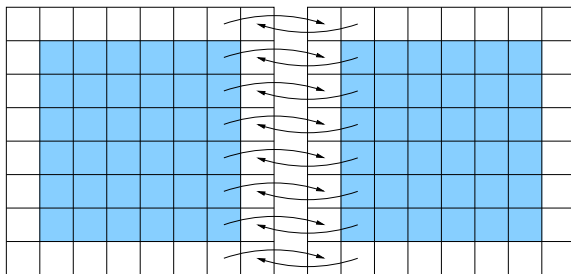
1. solve Riemann problems, compute wave speeds
2. compute maximum wave speed and infer global  $\Delta t$
3. update unknowns

# References & Literature

- LeVeque: *Finite Volume Methods for Hyperbolic Problems*, Cambridge University Press, 2002
- Toro: *Riemann Solvers and Numerical Methods for Fluid Dynamics: A Practical Introduction*, Springer, 2009
- Bale, LeVeque, Mitran, Rossmanith: *A wave propagation method for conservation laws and balance laws with spatially varying flux functions*, SIAM Journal on Scientific Computing 24 (3), 2003
- George: *Augmented Riemann solvers for the shallow water equations over variable topography with steady states and inundation*, Journal of Computational Physics 227 (6), 2008
- Breuer, Bader: *Teaching Parallel Programming Models on a Shallow-Water Code*, Proc. of the ISPDC 2012

# Part II

## Parallel Programming Patterns



*Reference: Mattson, Sanders, Massingill, Patterns for Parallel Programming. Addison-Wesley, 2005.*

# Finding Concurrency

## Common rule:

*Before you start parallelising your code, make sure the serial version is perfectly optimised!*

## Pro:

- parallelising a badly optimised serial algorithm leads to a badly optimised parallel algorithm
- **use an asymptotically optimal algorithm!**  
for large problems (that are worth being parallelised) asymptotics is crucial

## Contra:

- **exploit all available concurrency in your problem**  
(your optimised serial code might have unnecessary sequential parts)
- the fastest serial algorithm is not necessarily the fastest parallel algorithm



# Finding Concurrency – Task Decomposition

*Decompose your problem into tasks that can execute concurrently!*

**Consider “un-split” time stepping:**

$$\forall i,j: Q_{i,j}^{n+1} = Q_{i,j}^n - \frac{\Delta t}{\Delta x} \left( \mathcal{A}^+ \Delta Q_{i-1/2,j}^n + \mathcal{A}^- \Delta Q_{i+1/2,j}^n \right) - \frac{\Delta t}{\Delta y} \left( \mathcal{B}^+ \Delta Q_{i,j-1/2}^n + \mathcal{B}^- \Delta Q_{i,j+1/2}^n \right)$$

**Concurrent tasks:**

1. compute net updates (i.e., solve Riemann problems)  $\mathcal{A}^+ \Delta Q_{i-1/2,j}^n, \mathcal{B}^+ \Delta Q_{i,j-1/2}^n$  for all (vertical and horizontal) edges
  2. update quantities  $Q_{i,j}^{n+1}$  in all cells
- or:** for all cells, compute net updates (on local edges) and update quantities  $Q_{i,j}^{n+1}$  (requires two arrays for  $Q_{i,j}^n$  and  $Q_{i,j}^{n+1}$ , resp.)

# Finding Concurrency – Task Decomposition

*Decompose your problem into tasks that can execute concurrently!*

## Consider Dimensional Splitting:

$$Q_{i,j}^* = Q_{i,j}^n - \frac{\Delta t}{\Delta x} \left( \mathcal{A}^+ \Delta Q_{i-1/2,j}^n + \mathcal{A}^- \Delta Q_{i+1/2,j}^n \right)$$
$$Q_{i,j}^{n+1} = Q_{i,j}^* - \frac{\Delta t}{\Delta y} \left( \mathcal{B}^+ \Delta Q_{i,j-1/2}^* + \mathcal{B}^- \Delta Q_{i,j+1/2}^* \right).$$

## Concurrent tasks:

1. compute net updates on all vertical edges ( $\mathcal{A}^+ \Delta Q_{i-1/2,j}^n$ , etc.)
- 1a. update intermediate quantities  $Q_{i,j}^*$  in all cells
2. compute net updates on all horizontal edges ( $\mathcal{B}^+ \Delta Q_{i,j-1/2}^*$ , etc.)
- 2a. update quantities  $Q_{i,j}^{n+1}$  in all cells

# Finding Concurrency – Data Decomposition

*Decompose your data into units that can be operated on relatively independently!*

## Consider Dimensional Splitting:

$$Q_{i,j}^* = Q_{i,j}^n - \frac{\Delta t}{\Delta x} \left( \mathcal{A}^+ \Delta Q_{i-1/2,j}^n + \mathcal{A}^- \Delta Q_{i+1/2,j}^n \right)$$
$$Q_{i,j}^{n+1} = Q_{i,j}^* - \frac{\Delta t}{\Delta y} \left( \mathcal{B}^+ \Delta Q_{i,j-1/2}^* + \mathcal{B}^- \Delta Q_{i,j+1/2}^* \right).$$

## Data Decomposition:

1. computation of  $Q_{i,j}^*$ : distribute data row-wise, as computation is independent for different  $j$
2. update of  $Q_{i,j}^{n+1}$ : distribute data column-wise, as computation is independent for different  $i$

# Finding Concurrency – Data Decomposition

*Decompose your data into units that can operated on relatively independently!*

**Consider “un-split” time stepping:**

$$\forall i,j: Q_{i,j}^{n+1} = Q_{i,j}^n - \frac{\Delta t}{\Delta x} \left( \mathcal{A}^+ \Delta Q_{i-1/2,j}^n + \mathcal{A}^- \Delta Q_{i+1/2,j}^n \right) - \frac{\Delta t}{\Delta y} \left( \mathcal{B}^+ \Delta Q_{i,j-1/2}^n + \mathcal{B}^- \Delta Q_{i,j+1/2}^n \right)$$

**Concurrent tasks:**

- compute net updates requires left/right and top/down neighbours  
⇒ no “perfect” data decomposition possible
- partitioning of data will require extra care at **boundaries** of the partitions
- **and:** (seemingly trivial) do not decompose quantities in  $Q_{i,j}$

# Task and Data Decomposition – “Forces”

## Flexibility:

- be flexible enough to adapt to different implementation requirements
- for example: do not concentrate on a single parallel platform or programming model

## Efficiency:

- solution needs to scale efficiently with the size of the computer
- task and data decomposition need to provide enough tasks to keep all processing elements busy

## Simplicity:

- complex enough to solve the task, but simple enough to keep program maintainable

# Identifying Dependencies Between Tasks

## Group Tasks:

*Group your tasks to simplify the managing of dependencies*

## Order Tasks:

*Given a collection of tasks into logically related groups, order these task groups to satisfy constraints*

## Data Sharing:

*Given a data and task decomposition, how is data shared among the tasks?*

# Element Updates as Task Groups

Consider “un-split” time stepping:

$$\forall i,j: Q_{i,j}^{n+1} = Q_{i,j}^n - \frac{\Delta t}{\Delta x} \left( \mathcal{A}^+ \Delta Q_{i-1/2,j}^n + \mathcal{A}^- \Delta Q_{i+1/2,j}^n \right) - \frac{\Delta t}{\Delta y} \left( \mathcal{B}^+ \Delta Q_{i,j-1/2}^n + \mathcal{B}^- \Delta Q_{i,j+1/2}^n \right)$$

## Grouped Tasks:

- solve Riemann problems on the four cell edges
- update quantities  $Q_{i,j}$  from the net updates

## Data Dependencies:

- tasks access quantities  $Q_{i\pm 1,j\pm 1}^n$  of neighbour cells  
 $\Rightarrow$  two copies required for  $Q_{i,j}^n$  and  $Q_{i,j}^{n+1}$
- Riemann problem computed twice for each edge?

# Riemann Solves and Updates as Task Groups

**Consider Dimensional Splitting:**

$$Q_{i,j}^* = Q_{i,j}^n - \frac{\Delta t}{\Delta x} \left( \mathcal{A}^+ \Delta Q_{i-1/2,j}^n + \mathcal{A}^- \Delta Q_{i+1/2,j}^n \right)$$
$$Q_{i,j}^{n+1} = Q_{i,j}^* - \frac{\Delta t}{\Delta y} \left( \mathcal{B}^+ \Delta Q_{i,j-1/2}^* + \mathcal{B}^- \Delta Q_{i,j+1/2}^* \right).$$

**Separate Task Groups** (for each of the two steps):

- solve Riemann problems on all horizontal (vertical) cell edges
- update quantities  $Q_{i,j}$  of an entire column (row)

**Data Dependencies:**

- tasks access neighbours in either row or column direction
- requires extra storage to compute the net updates (results of the Riemann problems)



# Computation of the CFL Condition

Consider “un-split” time stepping:

$$\forall i,j: Q_{i,j}^{n+1} = Q_{i,j}^n - \frac{\Delta t}{\Delta x} \left( \mathcal{A}^+ \Delta Q_{i-1/2,j}^n + \mathcal{A}^- \Delta Q_{i+1/2,j}^n \right) - \frac{\Delta t}{\Delta y} \left( \mathcal{B}^+ \Delta Q_{i,j-1/2}^n + \mathcal{B}^- \Delta Q_{i,j+1/2}^n \right)$$

where  $\Delta t$  results from wave propagation speeds

**Sequential Order of Tasks:**

1. solve Riemann problems on the four cell edges  
(compute wave propagation speeds as partial results)
2. determine maximum wave speed for CFL condition  $\rightsquigarrow \Delta t$
3. update quantities  $Q_{i,j}$  from the net updates

# The Geometric Decomposition Pattern

*How can your algorithm be organized around a data structure that has been decomposed into concurrently updatable “chunks”?*

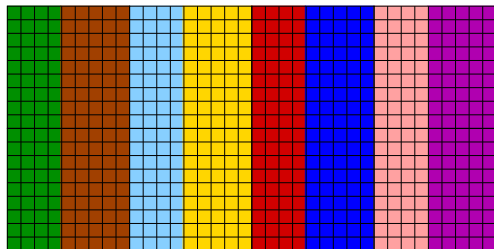
**Partitioning** (how to select your “chunks”):

- w.r.t. size, shape, etc. (“granularity” of parallelism)
- multiple levels of partitioning necessary?

**Organization** of parallel updates:

- need to access water height, momentum components and bathymetry from neighbour cells (possible in other partition)
- need to access net updates from neighbour partition?  
(alternative: compute on all involved partitions?)

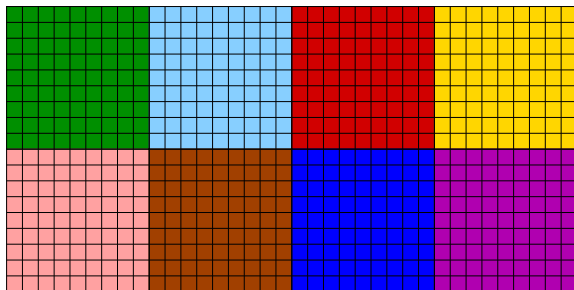
# 1D Domain Decomposition – Slice-Oriented



## Discussion:

- degenerates for large number of partitions: thin slices, lots of data exchange required at (long!) boundaries
- for dimensional splitting: slices match dependencies (vertical or horizontal) but alternating slices required for the two update steps

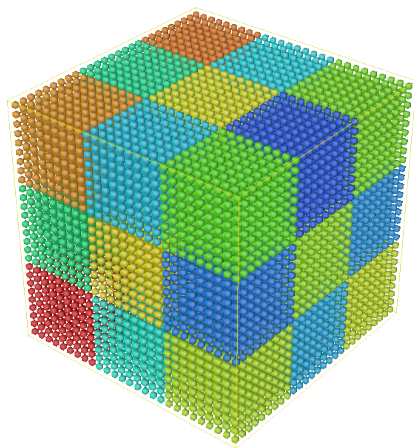
# 2D Domain Decomposition – Block-Oriented



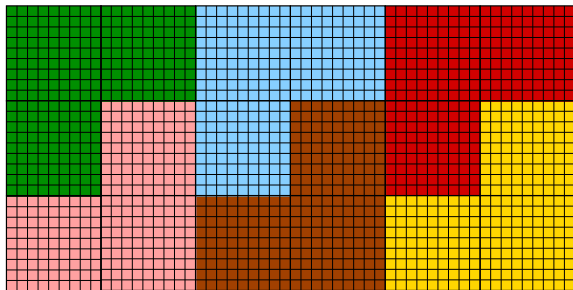
## Discussion:

- + length of domain boundaries (communication volume)
- fit arbitrary number of partitions to layout of boxes

# 3D Domain Decomposition – Cuboid-Oriented



# “Patches” Concept for Domain Decomposition

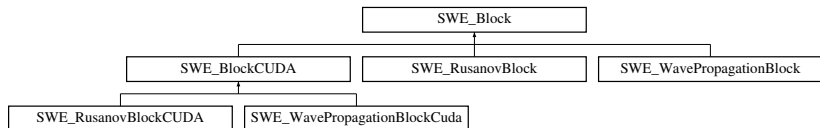


## Discussion:

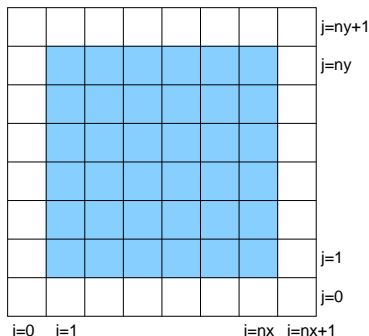
- + more fine-grain load distribution
- + “empty patches” improve representation of complicated domains
- overhead for additional, interior boundaries
- requires scheme to assign patches to processes

# Part III

## SWE Software Design



# Basic Structure: Cartesian Grid Block

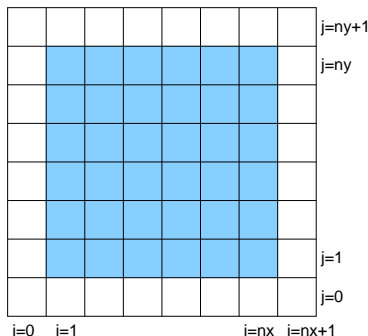


## Spatial Discretization:

- regular Cartesian meshes; later: allow multiple patches
- ghost layers to implement boundary conditions;  
connect multiple patches (complicated domains, parallelization)



# Basic Structure: Cartesian Grid Block



## Data Structure:

- arrays  $h$ ,  $h_u$ ,  $h_v$ , and  $b$  to hold water height, momentum components and bathymetry data
- “column major” layout:  $j$  the “faster running” index in  $h[i][j]$

# Main Loop – Euler Time-stepping

```
while( t < ... ) {  
    // set boundary conditions  
    splash.setGhostLayer();  
  
    // compute fluxes on each edge  
    splash.computeNumericalFluxes();  
  
    // set largest allowed time step:  
    dt = splash.getMaxTimestep();  
    t += dt;  
  
    // update unknowns in each cell  
    splash.updateUnknowns(dt);  
};
```

→ defines interface for abstract class **SWE\_Block**

# Set Ghost Layers – Boundary Conditions

Split into two methods:

- `setGhostLayer()`: interface function in `SWE_Block`, needs to be called by main loop
- `setBoundaryConditions()`: called by `setGhostLayer()`; sets “real” boundary conditions (WALL, OUTFLOW, etc.)

```
switch(boundary[BND_LEFT]) {  
  case WALL:  
  {  
    for(int j=1; j<=ny; j++) {  
      h[0][j] = h[1][j];    b[0][j] = b[1][j];  
      hu[0][j] = -hu[1][j]; hv[0][j] = hv[1][j];  
    };  
    break;  
  }  
  case OUTFLOW:  
  { /* ... */
```

(cmp. file `SWE_Block.cpp`)

# Compute Numerical Fluxes

main loop to compute net updates on **left/right edges**:

```
for(int i=1; i < nx+2; i++) {  
    for(int j=1; j < ny+1; j++) {  
        float maxEdgeSpeed;  
        wavePropagationSolver.computeNetUpdates(  
            h[i-1][j], h[i][j],  
            hu[i-1][j], hu[i][j],  
            b[i-1][j], b[i][j],  
            hNetUpdatesLeft[i-1][j-1], hNetUpdatesRight[i-1][j-1],  
            huNetUpdatesLeft[i-1][j-1], huNetUpdatesRight[i-1][j-1],  
            maxEdgeSpeed  
        );  
        maxWaveSpeed = std::max(maxWaveSpeed, maxEdgeSpeed);  
    }  
}
```

(cmp. file SWE\_WavePropagationBlock.cpp)

## Compute Numerical Fluxes (2)

main loop to compute net updates on **top/bottom edges**:

```
for(int i=1; i < nx+1; i++) {  
    for(int j=1; j < ny+2; j++) {  
        float maxEdgeSpeed;  
        wavePropagationSolver.computeNetUpdates(  
            h[i][j-1], h[i][j],  
            hv[i][j-1], hv[i][j],  
            b[i][j-1], b[i][j],  
            hNetUpdatesBelow[i-1][j-1], hNetUpdatesAbove[i-1][j-1],  
            hvNetUpdatesBelow[i-1][j-1], hvNetUpdatesAbove[i-1][j-1],  
            maxEdgeSpeed  
        );  
        maxWaveSpeed = std::max(maxWaveSpeed, maxEdgeSpeed);  
    }  
}
```

(cmp. file SWE\_WavePropagationBlock.cpp)

# Determine Maximum Time Step

- variable `maxWaveSpeed` holds maximum wave speed
- updated during computation of numerical fluxes in method `computeNumericalFluxes()`:  
$$\text{maxTimestep} = \text{std::min}(dx/\text{maxWaveSpeed}, dy/\text{maxWaveSpeed});$$
- simple “getter” method defined in class `SWE_Block`:  
**float** `getMaxTimestep()` { **return** `maxTimestep`; };
- hence: `getMaxTimestep()` for current time step should be called *after* `computeNumericalFluxes()`
- in general: in many situations, the maximum computation inhibits certain optimizations → fixed time step probably faster!

# Update Unknowns – Euler Time Stepping

```
for(int i=1; i < nx+1; i++) {  
    for(int j=1; j < ny+1; j++) {  
        h[i][j] -= dt/dx * (hNetUpdatesRight[i-1][j-1]  
                           + hNetUpdatesLeft[i][j-1] )  
        + dt/dy * (hNetUpdatesAbove[i-1][j-1]  
                  + hNetUpdatesBelow[i-1][j] );  
        hu[i][j] -= dt/dx * (huNetUpdatesRight[i-1][j-1]  
                             + huNetUpdatesLeft[i][j-1] );  
        hv[i][j] -= dt/dy * (hvNetUpdatesAbove[i-1][j-1]  
                             + hvNetUpdatesBelow[i-1][j] );  
        /* ... */  
    }  
}
```

(cmp. file SWE\_WavePropagationBlock.cpp)

# Goals for (Parallel) Implementation

## Spatial Discretization:

- allow different parallel programming models
- and also to switch between different numerical models

⇒ **class hierarchy of numerical vs. programming models**

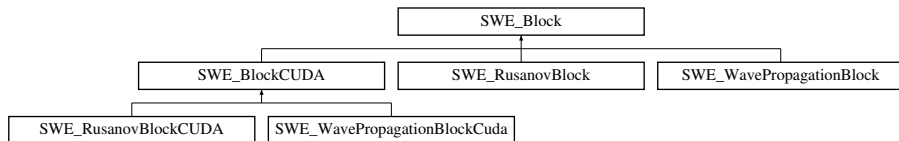
## Hybrid Parallelization:

- support two levels of parallelization  
(such as shared/distributed memory, CPU/GPU, etc.)
- coarse-grain parallelism across Cartesian grid patches
- fine-grain parallelism on patch-local operations

⇒ **separate fine-grain and coarse-grain parallelism**  
(plug&play principle)



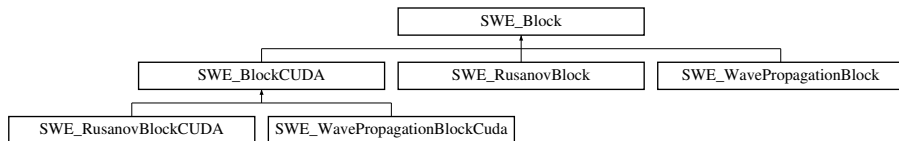
# SWE Class Design



## abstract class **SWE\_Block**:

- base class to hold data structures (arrays  $h$ ,  $h_u$ ,  $h_v$ ,  $b$ )
- manipulates ghost layers
- methods for initialization, writing output, etc.
- defines interface for main time-stepping loop:  
computeNumericalFluxes(), updateUnknowns(), ...

# SWE Class Design



## derived classes:

- for different model variants: `SWE_RusanovBlock`, `SWE_WavePropagationBlock`, ...
- for different programming models: `SWE_BlockCUDA`, `SWE_BlockArBB`, ...
- override `computeNumericalFluxes()`, `updateUnknowns()`, ...  
→ methods relevant for parallelization

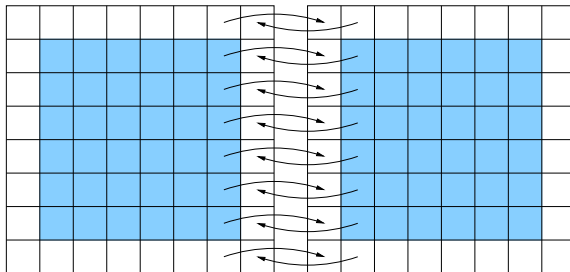
## Example: SWE\_WavePropagationBlockCUDA

```
class SWE_WavePropagationBlockCuda: public SWE_BlockCUDA {  
    /*— definition of member variables skipped —*/  
    public:  
        // compute a single time step (net-updates + update of the cells).  
        void simulateTimestep( float i_dT );  
        // simulate multiple time steps ( start and end time as parameters)  
        float simulate(float, float);  
        // compute the numerical fluxes (net-update formulation here).  
        void computeNumericalFluxes();  
        // compute the new cell values.  
        void updateUnknowns(const float i_deltaT);  
};
```

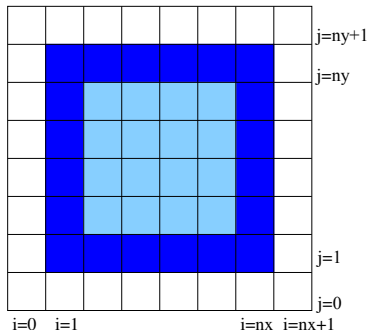
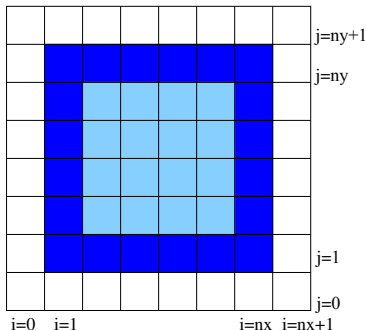
(in file SWE\_WavePropagationBlockCuda.hh)

# Part IV

## SWE Parallelisation



# Patches of Cartesian Grid Blocks



## Spatial Discretization:

- regular Cartesian meshes; allow multiple patches
- ghost and copy layers to implement boundary conditions, for more complicated domains, and for parallelization

# Loop-Based Parallelism within Patches

## Computing the Net Updates

- compute net updates on left/right edges:

```
for(int i=1; i < nx+2; i++) in parallel {
    for(int j=1; j < ny+1; j++) in parallel {
        float maxEdgeSpeed;
        fWaveComputeNetUpdates( 9.81,
                                h[i-1][j], h[i][j], hu[i-1][j], hu[i][j], /* ... */ );
    }
}
```

- compute net updates on top/bottom edges:

```
for(int i=1; i < nx+1; i++) in parallel {
    for(int j=1; j < ny+2; j++) in parallel {
        fWaveComputeNetUpdates( 9.81,
                                h[i][j-1], h[i][j], hv[i][j-1], hv[i][j], /* ... */ );
    }
} (function fWaveComputeNetUpdates() defined in file solver/FWaveCuda.h)
```

# Computing the Net Updates

## Options for Parallelism

### Parallelization of computations:

- compute all vertical edges in parallel
- compute all horizontal edges in parallel
- compute vertical & horizontal edges in parallel (task parallelism)

### Parallel access to memory:

- concurrent read to variables  $h$ ,  $h_u$ ,  $h_v$
- exclusive write access to net-update variables on edges

# Loop-Based Parallelism within Patches (2)

## Updating the Unknowns

- update unknowns from net updates on edges:

```

for(int i=1; i < nx+1; i++) in parallel {
    for(int j=1; j < ny+1; j++) in parallel {
        h[i][j] -= dt/dx * (hNetUpdatesRight[i-1][j-1]
                           + hNetUpdatesLeft[i][j-1])
               + dt/dy * (hNetUpdatesAbove[i-1][j-1]
                           + hNetUpdatesBelow[i-1][j]);
        hu[i][j] -= dt/dx * (huNetUpdatesRight[i-1][j-1]
                             + huNetUpdatesLeft[i][j-1]);

        /* ... */
    }
}

```



# Updating the Unknowns

## Options for Parallelism

### Parallelization of computations:

- compute all cells in parallel

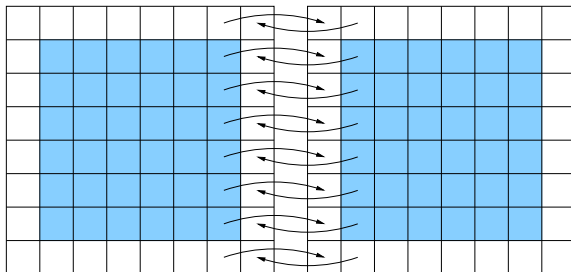
### Parallel access to memory:

- concurrent read to net-updates on edges
- exclusive write access to variables  $h$ ,  $h_u$ ,  $h_v$

### “Vectorization property”:

- exactly the same code for all cell!

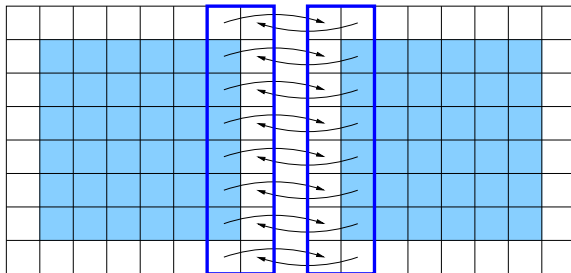
# Exchange of Values in Ghost/Copy Layers



## Straightforward Approach:

- boundary conditions OUTFLOW, WALL vs. CONNECT or PARALLEL
- disadvantage: method `setGhostLayer()` needs to be implemented for each derived class

## Exchange of Values in Ghost/Copy Layers (2)

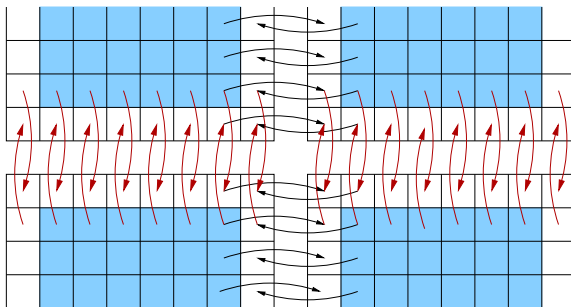


### Implemented via Proxy Objects:

- `grabGhostLayer()` to write into ghost layer
- `registerCopyLayer()` to read from copy layer
- both methods return a proxy object (class `SWE_Block1D`) that references one row/column of the grid

# Direct-Neighbour vs. “Diagonal” Communication

2-step scheme to exchange data of “diagonal” ghost cells:



- several “hops” replace diagonal communication
- slight increase of volume of communication (bandwidth), but reduces number of messages (latency)
- similar in 3D (26 neighbours → 6 neighbours!)

# MPI Parallelization

## – Exchange of Ghost/Copy Layers

```
SWE_Block1D* leftInflow = splash.grabGhostLayer(BND_LEFT);  
SWE_Block1D* leftOutflow = splash.registerCopyLayer(BND_LEFT);
```

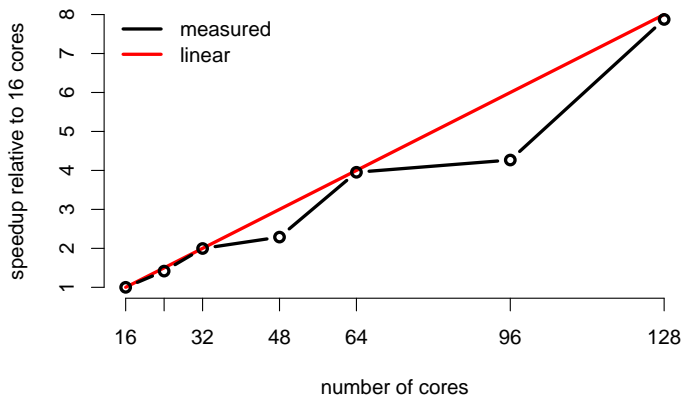
```
SWE_Block1D* rightInflow = splash.grabGhostLayer(BND_RIGHT);  
SWE_Block1D* rightOutflow = splash.registerCopyLayer(BND_RIGHT);
```

```
MPI_Sendrecv(leftOutflow->h.elemVector(), 1, MPI_COL, leftRank, 1,  
             rightInflow->h.elemVector(), 1, MPI_COL, rightRank, 1,  
             MPI_COMM_WORLD,&status);
```

```
MPI_Sendrecv(rightOutflow->h.elemVector(), 1, MPI_COL, rightRank, 4,  
             leftInflow->h.elemVector(), 1, MPI_COL, leftRank, 4,  
             MPI_COMM_WORLD,&status);
```

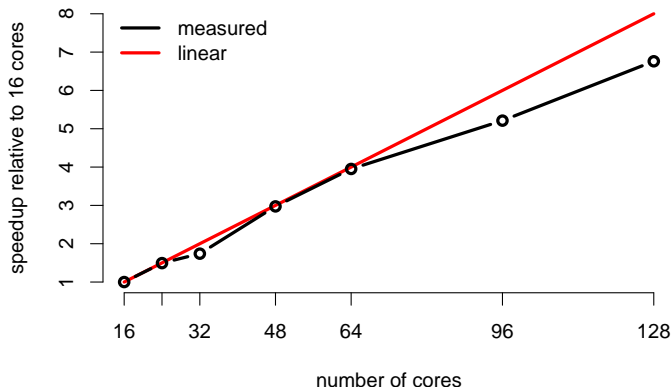
(cmp. file `examples/swe_mpi.cpp`)

# MPI – Some Speedups



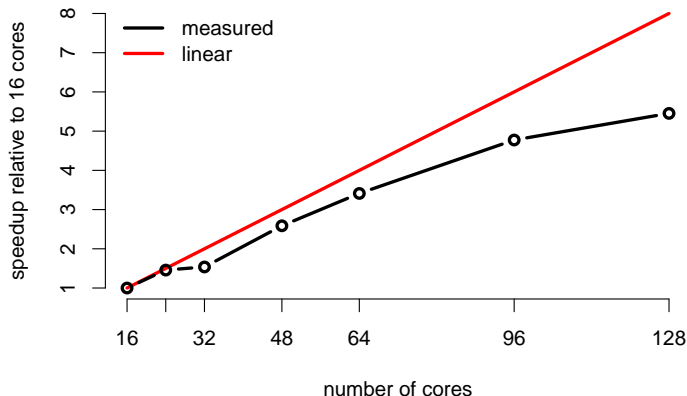
- 1 MPI process per core
- (expensive) augmented Riemann solvers

# Speedups for MPI/OpenMP



- 1 MPI process per node, 8 OpenMP threads (1 per core)
- straightforward OpenMP parallelization of for-loops

# Speedups for MPI/OpenMP



- 1 MPI process per node, 8 OpenMP threads (1 per core)
- hybrid f-Wave/aug. Riemann solver → poor load balancing



# Teaching Parallel Programming with SWE

## SWE in Lectures, Tutorials, Lab Courses:

- non-trivial example, but model & implementation easy to grasp
- allows different parallel programming models (MPI, OpenMP, CUDA, Intel TBB/ArBB, OpenCL, ...)
- prepared for hybrid parallelisation

## Some Extensions:

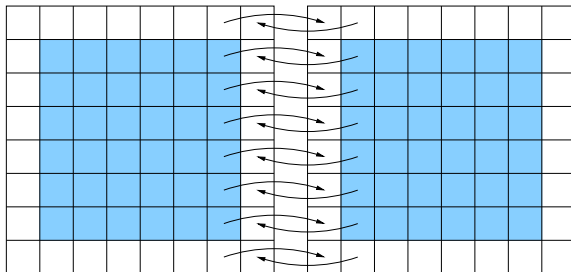
- ASAGI - parallel server for geoinformation (S. Rettenberger, Master's thesis)
- OpenGL real-time visualisation of results (T. Schnabel, student project; extended by S. Rettenberger)

→ <http://www5.in.tum.de/SWE/>

→ <https://github.com/TUM-I5>

# Part V

## Workshop – SWE Parallelisation



# MPI Communication Between Patches

## Extend sequential SWE program `swe_serial.cpp`:

- goal: one patch (SWE\_Block per MPI process)
- establish assignment of patches to MPI ranks (“who is my neighbour?”)
- implement exchange between ghost & copy cells (preferably via proxy objects)
- parallelize adaptive time step control
- produce speed-up graphs (strong and weak scaling)

## Possible extensions: (for the ambitious . . . )

- compare blocking vs. non-blocking communication
- try overlapping communication and computation
- allow multiple patches per MPI process

# Loop Parallelism in SWE Using OpenMP

## What should be done before starting with OpenMP?

- determine most time-consuming parts of your code (→ week 2)
- use option “guided auto-parallelism” of Intel compiler (→ welcome to try, but does not give many hints for SWE)

## Extend MPI-parallel SWE program towards MPI+OpenMP:

- what are the most time-consuming loops in SWE?
- ToDo: loop parallelism for these loops using **#pragma** ...
- test performance of OpenMP vs. MPI implementation

## Possible extensions: (for the ambitious ...)

- parallelize adaptive time step control (reduction)
- multiple-patch version: try OpenMP on patches