


Structured Grids and Space-Filling Curves

CSCS-FoMICS-USI Summer School on
Computer Simulations in Science and Engineering

Michael Bader



July 8–19, 2013



Part I



From Quadtrees to Space-Filling Curves

Quadrees to Describe Geometric Objects





- start with an initial square (covering the entire domain)
- recursive substructuring in four subsquares
- adaptive refinement possible
- terminate, if squares entirely within or outside domain

Storing a Quadtree – Sequentialisation



- sequentialise cell information according to **depth-first traversal**
- relative numbering of the child nodes determines sequential order
- here: leads to so-called **Morton order**

Morton Order



Relation to bit arithmetics:

- odd digits: position in vertical direction
- even digits: position in horizontal direction

Morton Order and Cantor's Mapping

Georg Cantor (1877):

$$0.\textcolor{blue}{0}\textcolor{red}{1}\textcolor{blue}{1}\textcolor{red}{1}\textcolor{blue}{0}\textcolor{red}{0}\textcolor{blue}{1} \dots \rightarrow \begin{pmatrix} 0.\textcolor{red}{0}\textcolor{red}{1}\textcolor{red}{1}\textcolor{red}{0} \dots \\ 0.\textcolor{blue}{1}\textcolor{blue}{1}\textcolor{blue}{0}\textcolor{blue}{1} \dots \end{pmatrix}$$


- **bijjective** mapping $[0, 1] \rightarrow [0, 1]^2$
- proved identical cardinality of $[0, 1]$ and $[0, 1]^2$
- provoked the question: is there a **continuous** mapping? (i.e. a curve)

Preserving Neighbourship for a 2D Octree

Requirements:



- consider a simple 4×4 -grid
- uniformly refined
- subsequently numbered cells should be neighbours in 2D

Leads to (more or less unique) numbering of children:



5	6	9	10
4	7	8	11
3	2	13	12
0	1	14	15

Preserving Neighbourship for a 2D Octree (2)



- adaptive refinement possible
- neighbours in sequential order remain neighbours in 2D
- here: similar to the concept of **Hilbert curves**

Open Questions

Algorithmics:

- How do we describe the sequential order algorithmically?
- What kind of operations are possible?
- Are there further “orderings” with the same or similar properties?

Applications:

- Can we quantify the “neighbour” property?
- In what applications can this property be useful?
- What further operations

Part II

Space-Filling Curves


Definition of a Space-filling Curve

Given a continuous, surjective mapping $f: \mathcal{I} \rightarrow \mathcal{Q} \subset \mathbb{R}^n$, then $f_*(\mathcal{I})$ is called a *space-filling curve*, if $|\mathcal{Q}| > 0$.

Comments:

- a *curve* is defined as the image $f_*(\mathcal{I})$ of a continuous mapping $f: \mathcal{I} \rightarrow \mathbb{R}^n$
- *surjective*: every element in \mathcal{Q} occurs as a value of f , i.e., $\mathcal{Q} = f_*(\mathcal{I})$
- $\mathcal{I} \subset \mathbb{R}$ and \mathcal{I} is compact, typically $\mathcal{I} = [0, 1]$
- if \mathcal{Q} is a smooth manifold, then there can be *no bijective* space-filling mapping $f: \mathcal{I} \rightarrow \mathcal{Q} \subset \mathbb{R}^n$ (theorem: E. Netto, 1879).


Example: Construction of the Hilbert curve



Iterations of the Hilbert curve:

- start with an iterative numbering of 4 subsquares
- combine four numbering patterns to obtain a twice-as-large pattern
- proceed with further iterations

Example: Construction of the Hilbert curve




Recursive construction of the iterations:

- split the quadratic domain into 4 congruent subsquares
- find a space-filling curve for each subdomain
- join the four subcurves in a suitable way

A Grammar for Describing the Hilbert Curve

Construction of the iterations of the Hilbert curve:



→ motivates a **Grammar** to generate the iterations

A Grammar for Describing the Hilbert Curve

- Non-terminal symbols: $\{H, A, B, C\}$, start symbol H
- terminal characters: $\{\uparrow, \downarrow, \leftarrow, \rightarrow\}$
- productions:

$$H \leftarrow A \uparrow H \rightarrow H \downarrow B$$

$$A \leftarrow H \rightarrow A \uparrow A \leftarrow C$$

$$B \leftarrow C \leftarrow B \downarrow B \rightarrow H$$

$$C \leftarrow B \downarrow C \leftarrow C \uparrow A$$

- replacement rule: in any word,
all non-terminals have to be replaced at the same time
 \rightarrow L-System (Lindenmayer)

\Rightarrow the arrows describe the **iterations of the Hilbert curve** in “turtle graphics”

Definition of the Hilbert Curve's Mapping

Definition: (Hilbert curve)

- each parameter $t \in \mathcal{I} := [0, 1]$ is contained in a sequence of intervals

$$\mathcal{I} \supset [a_1, b_1] \supset \dots \supset [a_n, b_n] \supset \dots,$$

where each interval result from a division-by-four of the previous interval.

- each such sequence of intervals can be uniquely mapped to a corresponding sequence of 2D intervals (subsquares)
- the 2D sequence of intervals converges to a unique point q in $q \in \mathcal{Q} := [0, 1] \times [0, 1]$ – q is defined as $h(t)$.

Theorem


$h : \mathcal{I} \rightarrow \mathcal{Q}$ defines a space-filling curve, the Hilbert curve.

Claim: h defines a Space-filling Curve


We need to prove:

- h is a mapping, i.e. each $t \in \mathcal{I}$ has a *unique* function value $h(t) \rightarrow \text{OK}$, if $h(t)$ is independent of the choice of the sequence of intervals (proof skipped)
- $h: \mathcal{I} \rightarrow \mathcal{Q}$ is *surjective*:
 - for each point $q \in \mathcal{Q}$, we can construct an appropriate sequence of 2D-intervals
 - the 2D sequence corresponds in a unique way to a sequence of intervals in \mathcal{I} – this sequence defines an original value of q
 \Rightarrow every $q \in \mathcal{Q}$ occurs as an image point.
- h is *continuous* \rightarrow see proof of *Hölder continuity*

3D Hilbert Curves – Iterations



1st iteration



2nd iteration

Part III

Parallelisation Using Space-Filling Curves

Generic Space-filling Heuristic


Bartholdi & Platzman (1988):

1. Transform the problem in the unit square, via a space-filling curve, to a problem on the unit interval
2. Solve the (easier) problem on the unit interval

For parallelisation: strategy to determine partitions


1. use a space-filling curve to generate a sequential order on the grid cells
2. do a 1D partitioning on the list of cells
(cut into equal-sized pieces, or similar)

Hilbert-Curve Partitions on a Cartesian Grid



- Hilbert curve splits vertices into right/left (red/green) set
- Hilbert order traversal provides boundary vertices in sequential order


Example: Hilbert-Curve Partitions on Quadrees



- here: with ghost cells
(processed in identical order in both partitions)

Recall: Grammar to Describe the Hilbert Curve

Construction of the iterations of the Hilbert curve:



Can this grammar be used to generate *adaptive* Hilbert orders?

A Grammar for Hilbert Orders on Quadrees

- Non-terminal symbols: $\{H, A, B, C\}$, start symbol H
- terminal characters: $\{\uparrow, \downarrow, \leftarrow, \rightarrow, (,)\}$
- productions:

$$H \leftarrow (A \uparrow H \rightarrow H \downarrow B)$$

$$A \leftarrow (H \rightarrow A \uparrow A \leftarrow C)$$

$$B \leftarrow (C \leftarrow B \downarrow B \rightarrow H)$$

$$C \leftarrow (B \downarrow C \leftarrow C \uparrow A)$$

- ⇒ arrows describe the iterations of the Hilbert curve in “turtle graphics”
- ⇒ terminals (and) mark change of levels: “up” and “down”

Hölder Continuity

A function $f: \mathcal{I} \rightarrow \mathbb{R}^n$ is (uniformly) *continuous*, if
for each $\epsilon > 0$ there is a $\delta > 0$, such that:
for all $t_1, t_2 \in \mathcal{I}$ with $|t_1 - t_2| < \delta$,
the image points have a distance of $\|f(t_1) - f(t_2)\|_2 < \epsilon$

Hölder Continuity:

f is called *Hölder continuous with exponent r* on \mathcal{I} ,
if a constant $C > 0$ exists, such that for all $t_1, t_2 \in \mathcal{I}$:

$$\|f(t_1) - f(t_2)\|_2 \leq C |t_1 - t_2|^r$$

- case $r = 1$ is equivalent to Lipschitz continuity
- Hölder continuity implies uniform continuity

Hölder Continuity and Parallelisation

$$\|f(t_1) - f(t_2)\|_2 \leq C |t_1 - t_2|^r$$

Interpretation:

- $\|f(t_1) - f(t_2)\|_2$ is the distance of the image points
- $|t_1 - t_2|$ is the distance of the indices
- also: $|t_1 - t_2|$ is the area of the respective space-filling-curve partition
- hence: relation between volume (number of grid cells/points) and extent (e.g. radius) of a partition

⇒ Hölder continuity gives a quantitative estimate for
compactness of partitions

Hölder Continuity of the Hilbert Curve

Proof:



- given $t_1, t_2 \in \mathcal{I}$; choose n , such that $4^{-(n+1)} < |t_1 - t_2| < 4^{-n}$
- 4^{-n} is interval length for the n -th iteration
 $\Rightarrow [t_1, t_2]$ overlaps at most two neighbouring(!) intervals.
- due to construction of the Hilbert curve, $h(t_1)$ and $h(t_2)$ are in neighbouring subsquares with face length 2^{-n} .
- these two subsquares build a rectangle with a diagonal of length $2^{-n} \cdot \sqrt{5}$; therefore: $\|h(t_1) - h(t_2)\|_2 \leq 2^{-n} \sqrt{5}$
- as $4^{-(n+1)} < |t_1 - t_2|$, we have $2 \cdot 2^{-n} < \sqrt{|t_1 - t_2|}$

\Rightarrow result: $\|h(t_1) - h(t_2)\|_2 \leq \frac{1}{2} \sqrt{5} |t_1 - t_2|^{1/2}$

Part IV



Outlook: Parallelisation with Space-Filling Curves and Refinement Trees

Hilbert-Order Bitstream-Encoding of a Quadtree



1 1 0 1 1 1 1 1 0 0 0 0 0 0 1 1 1 1 0 0 1 1 0 0 1 1 0 1 0 0 1 1

Refinement-Tree Encoding of a Quadtree





172 44 0 16 8 4 1 0 0 0 0 0 16 48 12 16 16 0 24 12 8 0 0 52 12 12 12 12

Refinement-Tree Encoding of a Quadtree (2)

REFTREE algorithm for partitioning:

- attributed quadtree with number of leaves/nodes of the subtree for each node
- allows to determine whether a certain node/subtree may be skipped by the current partition (if index of first & last leave/node are given)
- disadvantage of data structure: required information spread across several locations in the stream
⇒ may be fixed by modified depth-first order

Refinement-Tree Encoding with Modified Depth-First Order




172 44 48 24 52 0 16 8 16 4 0 0 0 1 0 0 0 12 16 16 0 12 8 0 0 12 12 12 12

(numbers in the tree represent position of resp. node information in the stream)

Parallelisation vs. Partitioning with SFC

Besides partitioning, parallelisation has to deal with:


- data exchange between partitions:
 - unknowns on separator between partitions
 - synchronize refinement status at partition boundaries
- may exploit “stack property” of Hilbert curves




How to Determine Left and Right

Turtle Grammars Revisited


H:




B:




L:




R:



E:




T:



Literature/References

<http://www.space-filling-curves.org/>



... and the references therein :-)