


Hybrid Architectures – Why Should I Bother?



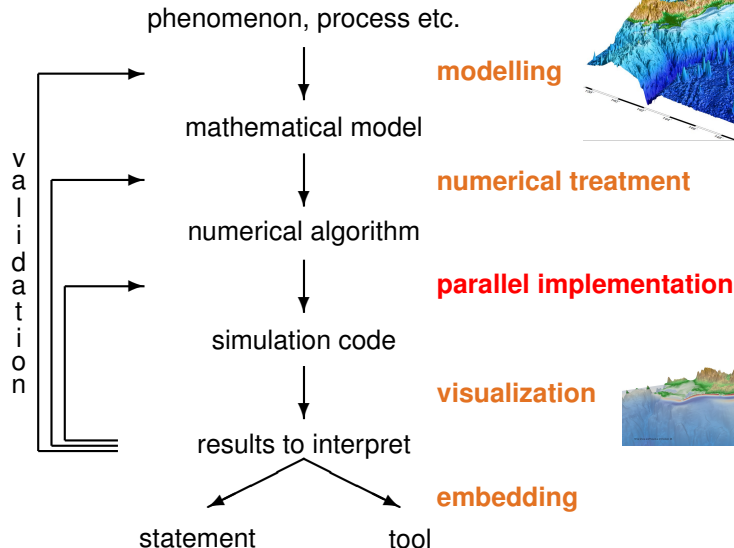
CSCS-FoMICS-USI Summer School on
Computer Simulations in Science and Engineering

Michael Bader

July 8–19, 2013



The Simulation Pipeline



Parallel Computing – “Faster, Bigger, More”

Why parallel high performance computing:

- **Response time:** compute a problem in $\frac{1}{p}$ time
 - speed up engineering processes
 - real-time simulations (tsunami warning?)
- **Problem size:** compute a p -times bigger problem
 - simulation of large-/multi-scale phenomena
 - maximal problem size that “fits into the machine”
 - validation of smaller, “operational” models
- **Throughput:** compute p problems at once
 - case and parameter studies, statistical risk scenarios, etc. (hazard maps, data base for tsunami warning, ...)
 - massively distributed computing (SETI@home, e.g.)

Part I

High Performance Computing in CSE – Past(?) and Present Trends

The Seven Dwarfs of HPC

“**dwarfs**” = key algorithmic kernels in many scientific computing applications

P. Colella (LBNL), 2004:

1. dense linear algebra
2. sparse linear algebra
3. spectral methods
4. N-body methods
5. **structured grids**
6. **unstructured grids**
7. Monte Carlo



Tsunami & storm-surge simulation:

- usually PDE solvers on structured or unstructured meshes
- **SWE**: a simple shallow water solver on Cartesian grids

Computational Science Demands a New Paradigm

Computational simulation must meet three challenges to become a mature partner of theory and experiment

(Post & Votta, 2005)

1. performance challenge

- exponential growth of performance, massively parallel architectures

2. programming challenge

- new (parallel) programming models

3. prediction challenge


- careful verification and validation of codes;
towards **reproducible** simulation experiments

Four Horizons for Enhancing the Performance ...

... of Parallel Simulations Based on Partial Differential Equations
(David Keyes, 2000)

1. Expanded Number of Processors
 - in 2000: 1000 cores; in 2010: 200,000 cores
2. More Efficient Use of Faster Processors
 - PDF working-sets, cache efficiency
3. More Architecture-Friendly Algorithms
 - improve temporal/spatial locality
4. Algorithms Delivering More “Science per Flop”
 - adaptivity (in space and time), higher-order methods, fast solvers

Performance Development in Supercomputing



(source: www.top500.org)

Top 500 (www.top500.org) – June 2013

Rank	Site	System	Cores	Rmax (TFlop/s)	Rpeak (TFlop/s)	Power (kW)
1	National University of Defense Technology China	Tianhe-2 (MilkyWay-2) - TH-IVB-FEP Cluster, Intel Xeon E5-2692 12C 2.200GHz, TH Express-2, Intel Xeon Phi 31S1P NUDT	3120000	33862.7	54902.4	17808
2	DOE/SC/Oak Ridge National Laboratory United States	Titan - Cray XK7 , Opteron 6274 16C 2.200GHz, Cray Gemini interconnect, NVIDIA K20x Cray Inc.	560640	17590.0	27112.5	8209
3	DOE/NNSA/LLNL United States	Sequoia - BlueGene/Q, Power BQC 16C 1.60 GHz, Custom IBM	1572864	17173.2	20132.7	7890
4	RIKEN Advanced Institute for Computational Science (AICS) Japan	K computer , SPARC64 VIIIfx 2.0GHz, Tofu interconnect Fujitsu	705024	10510.0	11280.4	12660
5	DOE/SC/Argonne National Laboratory United States	Mira - BlueGene/Q, Power BQC 16C 1.60GHz, Custom IBM	786432	8586.6	10066.3	3945
6	Texas Advanced Computing Center/Univ. of Texas United States	Stampede - PowerEdge C8220, Xeon E5-2680 8C 2.700GHz, Infiniband FDR, Intel Xeon Phi SE10P Dell	462462	5168.1	8520.1	4510

Top 500 Spotlights – Tianhe-2 and K Computer

Tianhe-2/MilkyWay-2 → Intel Xeon Phi (NUDT)

- 3.1 mio cores(!) – Intel Ivy Bridge and Xeon Phi
- Linpack benchmark: 33.8 PFlop/s
- ≈ 17 MW power(!!)
- Knights Corner / **Intel Xeon Phi** / Intel MIC as accelerator
- 61 cores, roughly 1.1–1.3 GHz

Titan → Cray XK7, NVIDIA K20x (ORNL)

- 18,688 compute nodes; 300,000 Opteron cores
- 18,688 **NVIDIA Tesla K20 GPUs**
- Linpack benchmark: 17.6 PFlop/s
- ≈ 8.2 MW power

Top 500 Spotlights – Sequoia and K Computer


Sequoia → IBM BlueGene/Q (LLNL)

- 98,304 compute nodes; 1.6 mio cores
- Linpack benchmark: 17.1 PFlop/s
- \approx 8 MW power

K Computer → SPARC64 (RIKEN, Kobe)

- 88,128 processors; 705,024 cores
- Linpack benchmark: 10.51 PFlop/s
- \approx 12 MW power
- SPARC64 VIIIfx 2.0GHz (8-core CPU)

Performance Development in Supercomputing



(source: www.top500.org)

International Exascale Software Project Roadmap

Towards an Exa-Flop/s Platform in 2018 (www.exascale.org):

1. technology trends

- concurrency, reliability, power consumption, ...
- blueprint of an exascale system: 10-billion-way concurrency, 100 million to 1 billion cores, 10-to-100-way concurrency per core, hundreds of cores per die, ...

2. science trends

- climate, high-energy physics, nuclear physics, fusion energy sciences, materials science and chemistry, ...

3. X-stack (software stack for exascale)

- energy, resiliency, heterogeneity, I/O and memory

4. Politico-economic trends

- exascale systems run by government labs, used by CSE scientists

Exascale Roadmap

“Aggressively Designed Strawman Architecture”

Level	What	Perform.	Power	RAM
FPU	FPU, regs., instr.-memory	1.5 Gflops	30mW	
Core	4 FPUs, L1	6 Gflops	141mW	
Proc. Chip	742 cores , L2/L3, Interconn.	4.5 Tflops	214W	
Node	Proc. chip, DRAM	4.5 Tflops	230W	16 GB
Group	12 proc. chips, routers	54 Tflops	3.5KW	192 GB
rack	32 groups	1.7 Pflops	116KW	6.1 TB
System	583 racks	1 Eflops	67.7MW	3.6 PB

approx. 285,000 cores per rack; 166 mio cores in total

Source:


ExaScale Computing Study: Technology Challenges in Achieving Exascale Systems

Exascale Roadmap – Should You Bother?

Your department's compute cluster in 5 years?

- a **Petaflop System!**
- “one rack of the Exaflop system”
→ using the same/similar hardware
- extrapolated example machine:
 - peak performance: 1.7 PFlop/s
 - 6 TB RAM, 60 GB cache memory
 - “total concurrency”: $1.1 \cdot 10^6$
 - number of cores: 280,000
 - number of chips: 384

Source: ExaScale Software Study: Software Challenges in Extreme Scale Systems



Your Department's PetaFlop/s Cluster in 5 Years?

Tianhe-1A (Tianjin, China; Top500 # 10)

- 14,336 Xeon X5670 CPUs
- **7,168 Nvidia Tesla M2050 GPUs**
- Linpack benchmark: ≈ 2.6 PFlop/s
- ≈ 4 MW power

Stampede (Intel, Top500 # 6)


- 102,400 cores (incl. Xeon Phi: MIC/"many integrated cores")
- Linpack benchmark: ≈ 5 PFlop/s
- Knights Corner / **Intel Xeon Phi** / Intel MIC as accelerator
- 61 cores, roughly 1.1–1.3 GHz
- wider vector FP units: 64 bytes (i.e., 16 floats, 8 doubles)
- ≈ 4.5 MW power

Free Lunch is Over^(*)

... actually already over for quite some time!

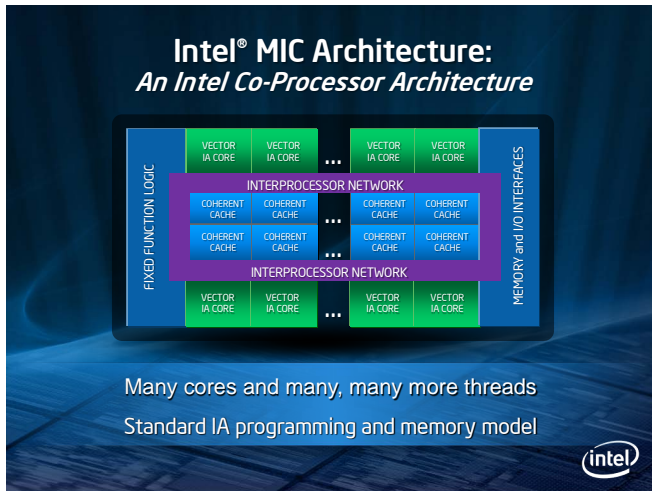
Speedup of your software can only come from parallelism:

- clock speed of CPU has stalled
- instruction-level parallelism per core has stalled
- number of cores is growing
- size of vector units is growing



(*) Quote and image taken from: H. Sutter, *The Free Lunch Is Over: A Fundamental Turn Toward Concurrency in Software*, Dr. Dobbs's Journal 30(3), March 2005.

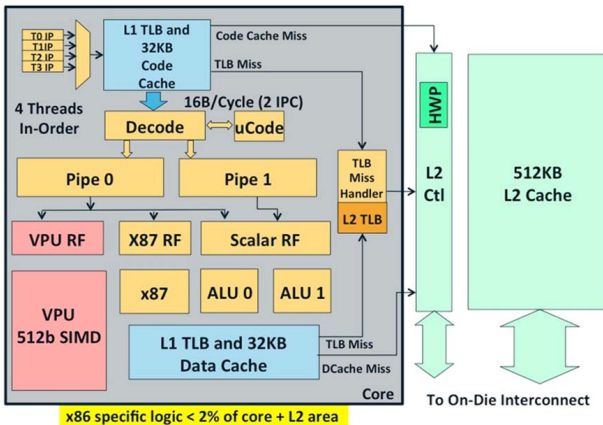
Manycore CPU – Intel MIC Architecture



(source: Intel/K. Skaugen – SC'10 keynote presentation)


Manycore CPU – Intel MIC Architecture (2)

Diagram of a Knights Corner core:




(source: An Overview of Programming for Intel Xeon processors and Intel Xeon Phi coprocessors)

GPGPU – NVIDIA Fermi

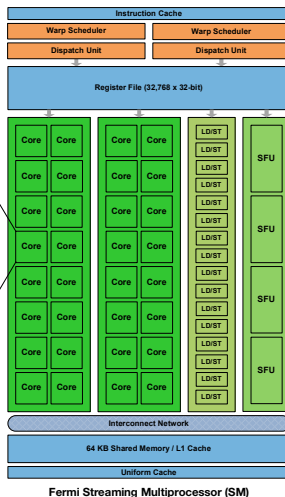


(source: NVIDIA – Fermi Whitepaper)

GPGPU – NVIDIA Fermi (2)



(source: NVIDIA – Fermi Whitepaper)




GPGPU – NVIDIA Fermi (3)

General Purpose Graphics Processing Unit:

- 512 CUDA cores
- improved double precision performance
- shared vs. global memory
- new: L1 und L2 cache (768 KB)
- trend from GPU towards CPU?

Fermi Memory Hierarchy



Parallel Computing Paradigms

Not exactly sure how the hardware will look like ...
(CPU-style, GPU-style, something new?)

However: **massively parallel programming** required

- revival of vector computing
→ several/many FPU's performing the same operation
- hybrid/heterogenous architectures
→ different kind of cores; dedicated accelerator hardware
- different access to memory
→ cache and cache coherency
→ small amount of memory per core
- our concern in this course: **data parallelism**
→ vectorisation (and a look into GPU computing)